

Reducer Results

August 24, 2021

[]:

```
[2]: #!/usr/bin/env python3
      """A more advanced Reducer, using Python iterators and generators."""

      from itertools import groupby
      from operator import itemgetter
      import sys
      import numpy as np
      import time
      import pandas as pd
      from IPython.display import display
      from datetime import datetime, timedelta
      from sklearn.cluster import MiniBatchKMeans, KMeans
      from pandarallel import pandarallel # use: pip install pandarallel
      from sklearn.model_selection import train_test_split
      from sklearn.linear_model import LinearRegression
      from xgboost import XGBRegressor # use: conda install xgboost
      from sklearn.ensemble import RandomForestRegressor
      from sklearn.metrics import r2_score
      from sklearn.metrics import mean_squared_error
      import math
      from prettytable import PrettyTable
      import gc
      import fileinput
      start_time = time.time()
      i = 0

      def model_evaluation(algorithem_name, X_Test, y_pred, y_true):

          # R2 and Adjasted R2
          r2 = r2_score(y_true, y_pred)
          adj_r2 = 1-(1-r2)*((len(X_Test)-1)/(len(X_Test)-X_Test.shape[1]-1))
          # MSE and RMSE
          mse = mean_squared_error(y_true, y_pred)
          rmse = math.sqrt(mse)
```

```

# print in table
table = PrettyTable()
print('\n' + algorithm_name + ':')
table.field_names = ['Experiment', 'Value']
table.add_row(['R2', r2])
table.add_row(['Adjusted R2', adj_r2])
table.add_row(['MSE', mse])
table.add_row(['RMSE', rmse])
return table

def measure_time(label = ''):
    global start_time
    global i
    print('{}\t took {}s'.format(label, round(time.time() - start_time)))
    start_time = time.time()
    i+=1

def read_mapper_output(file):
    for line in file:
        yield line

def main(separator='\t'):

    df_train = []
    df_test = []
    names=['pickup_time', 'pickup_longitude', 'pickup_latitude']
    c=['tpep_pickup_datetime',
        'pickup_longitude',
        'pickup_latitude']
    measure_time('initialisation')

    # input comes from STDIN (standard input)
    gc.disable()

    with fileinput.input(files=('/home/Aziz/MapperOutput_2015-01.csv',
                                '/home/Aziz/MapperOutput_2015-02.csv',
                                '/home/Aziz/MapperOutput_2015-03.csv',
                                '/home/Aziz/MapperOutput_2016-01.csv',
                                '/home/Aziz/MapperOutput_2016-02.csv',
                                '/home/Aziz/MapperOutput_2016-03.csv'
                                )) as f:

        for line in f:
            #data = read_mapper_output(sys.stdin)
            for row in f:

```

```

        #print(row)
        key, row = row.split(separator)
        row = row.split(',')
        if key=='train':
            df_train.append(row)
        elif key=='test':
            df_test.append(row)

#     data = read_mapper_output(sys.stdin)
#     for row in data:
#         print(row)
#         key, row = row.split(separator)
#         row = row.split(',')
#         if key=='train':
#             df_train.append(row)
#         elif key=='test':
#             df_test.append(row)
#     gc.enable()
measure_time('split rows') # measuring time for debugging purposes

# converting rows of data into pandas dataframes
df_train = pd.DataFrame(df_train, columns=names)
df_test = pd.DataFrame(df_test, columns=names)
# for debug only
print('No. rows in training dataset: %d' % df_train.shape[0])
print('No. rows in test dataset: %d' % df_test.shape[0])
measure_time('convert rows to dataframes') # printing time for debugging
→only

## GEOGRAPHICAL SEGMENTATION BY CLUSTERING
#Clustering pickups, Getting clusters
coord = df_train[["pickup_latitude", "pickup_longitude"]].values
regions = MiniBatchKMeans(n_clusters = 30, batch_size = 10000).fit(coord)
measure_time('training cluster model') # printing time for debugging only

#predicting clusters
df_train["pickup_cluster"] = regions.predict(df_train[["pickup_latitude",
→"pickup_longitude"]])
df_test["pickup_cluster"] = regions.predict(df_test[["pickup_latitude",
→"pickup_longitude"]])

```

```

measure_time('predicting clusters') # printing time for debugging only

# Replacing mins and sec with 0
pandarallel.initialize()
df_train['pickup_time'] = df_train.pickup_time.parallel_apply(lambda x : pd.
→to_datetime(x).replace(minute=0, second=0))
df_test['pickup_time'] = df_test.pickup_time.parallel_apply(lambda x : pd.
→to_datetime(x).replace(minute=0, second=0))
measure_time('reformat time') # printing time for debugging only

## CALCULATING TAXI-DEMAND
# Group by Cluster_id and time
df_train_2 = df_train.groupby(['pickup_time', 'pickup_cluster']).size().
→reset_index(name='count')
df_test_2 = df_test.groupby(['pickup_time', 'pickup_cluster']).size().
→reset_index(name='count')
measure_time('grouping and calculate pickups') # printing time for
→debugging only

# Converting pickup counts to demand percentage
df_train_2['count'] = df_train_2['count'].parallel_apply(lambda x : (x /
→df_train_2['count'].max()))
df_test_2['count'] = df_test_2['count'].parallel_apply(lambda x : (x /
→df_test_2['count'].max()))
measure_time('pickups to percentage') # printing time for debugging only

# Getting month, days, hours, day of week
df_train_2['month'] = pd.DatetimeIndex(df_train_2['pickup_time']).month
df_train_2['day'] = pd.DatetimeIndex(df_train_2['pickup_time']).day
df_train_2['dayofweek'] = pd.DatetimeIndex(df_train_2['pickup_time']).
→dayofweek
df_train_2['hour'] = pd.DatetimeIndex(df_train_2['pickup_time']).hour

df_test_2['month'] = pd.DatetimeIndex(df_test_2['pickup_time']).month
df_test_2['day'] = pd.DatetimeIndex(df_test_2['pickup_time']).day
df_test_2['dayofweek'] = pd.DatetimeIndex(df_test_2['pickup_time']).
→dayofweek
df_test_2['hour'] = pd.DatetimeIndex(df_test_2['pickup_time']).hour

```

```

measure_time('reformat time') # printing time for debugging only

# X and y for training
X_train = df_train_2[['pickup_cluster', 'month', 'day', 'hour',
↳ 'dayofweek']]
y_train = df_train_2['count']

# X and y for testing
X_test = df_test_2[['pickup_cluster', 'month', 'day', 'hour', 'dayofweek']]
y_test = df_test_2['count']
measure_time('splitting to train and test') # printing time for debugging
↳ only

# split training data into training and validation data
#X_train, X_validation, y_train, y_validation = train_test_split(X_train,
↳ y_train,
#
test_size=0.33,
↳ random_state=42)
#measure_time('splitting train and validation') # printing time for
↳ debugging only

# Linear regression
LReg = LinearRegression()
LReg.fit(X_train, y_train)
LReg_y_pred = LReg.predict(X_test)

# RandomForest regression
RFRegr = RandomForestRegressor()
RFRegr.fit(X_train, y_train)
RFRegr_y_pred = RFRegr.predict(X_test)

# XGB regression
GBRegr = XGBRegressor(n_estimators=1000, max_depth=7, eta=0.1, subsample=0.
↳ 7, colsample_bytree=0.8)
GBRegr.fit(X_train, y_train)
GBRegr_y_pred = GBRegr.predict(X_test)
measure_time('training models') # printing time for debugging only

# evaluation

```

```

    #print(model_evaluation('y True',X_Test=X_test, y_pred=y_test,
↪y_true=y_test))
    print(model_evaluation('Linear Regression',X_Test=X_test,
↪y_pred=LReg_y_pred, y_true=y_test))
    print(model_evaluation('RandomForest',X_Test=X_test, y_pred=RFRegr_y_pred,
↪y_true=y_test))
    print(model_evaluation('XGB',X_Test=X_test, y_pred=GBRegr_y_pred,
↪y_true=y_test))
    measure_time('printing results') # printing time for debugging only

    #print("--- %s seconds --- 0" % (time.time() - start_time))
    #display(X_train)
    #display(y_train)
    #display(X_validation)
    #display(y_validation)
    #display(X_test)
    #display(y_test)

if __name__ == "__main__":
    main()

```

```

initialisation    took 0s
split rows        took 97s
No. rows in training dataset:  37561189
No. rows in test dataset:  33722144
convert rows to dataframes      took 67s
training cluster model    took 36s
predicting clusters        took 69s
INFO: Pandarallel will run on 24 workers.
INFO: Pandarallel will use Memory file system to transfer data between the main
process and workers.
reformat time    took 597s
grouping and calculate pickups    took 4s
pickups to percentage    took 24s
reformat time    took 0s
splitting to train and test      took 0s

/opt/conda/anaconda/lib/python3.7/site-packages/sklearn/ensemble/forest.py:245:
FutureWarning: The default value of n_estimators will change from 10 in version
0.20 to 100 in 0.22.
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)

training models    took 13s

```

Linear Regression:

Experiment	Value
R2	0.12358378989809549
Adjusted R2	0.12351541283175582
MSE	0.022728725137356843
RMSE	0.15076048931121458

RandomForest:

Experiment	Value
R2	0.9121364300505471
Adjusted R2	0.9121295750276915
MSE	0.0022786284734918785
RMSE	0.04773498165383411

XGB:

Experiment	Value
R2	0.9190748065302935
Adjusted R2	0.919068492832237
MSE	0.002098690619661767
RMSE	0.045811468211156114

printing results took 0s

```
[48]: #sys.stdin = open('/home/Aziz/MapperOutput.txt', 'r')
```